

Graph Enhanced Transformer for Semi-Supervised Duplicate Bug Report Detection

1 Kerem Bayramoğlu*

2 kerem.bayramoglu@bilkent.edu.tr
3 Bilkent University
4 Dept. of Computer Engineering
5 Çankaya, Ankara, Türkiye

6 Hüseyin Karaca

7 huseyin.karaca@bilkent.edu.tr
8 Bilkent University
9 Dept. of Electrical and Electronics
10 Engineering
11 Çankaya, Ankara, Türkiye

12 Emirhan Koç

13 emirhan.koc@bilkent.edu.tr
14 Bilkent University
15 Dept. of Electrical and Electronics
16 Engineering
17 Çankaya, Ankara, Türkiye

18 Hasan Erkin Ünlü

19 erkin.unlu@bilkent.edu.tr
Bilkent University
Dept. of Electrical and Electronics
Engineering
Çankaya, Ankara, Türkiye

20 Rabia Ela Ünlü

21 ela.unlu@bilkent.edu.tr
Bilkent University
Dept. of Electrical and Electronics
Engineering
Çankaya, Ankara, Türkiye

ABSTRACT

Context. Duplicate bug report detection aims to identify bug reports that describe the same underlying problem or can be resolved by the same fix. In large-scale repositories, duplicates increase triage workload and waste developer time. While transformer-based sentence encoders improve semantic matching compared to classical IR baselines, they rely on labeled duplicate pairs, which are typically scarce in real-world repositories. Moreover, exhaustive pairwise training is computationally infeasible and cannot effectively incorporate the full pool of unlabeled reports.

Method. We propose a semi-supervised Graph-Enhanced Transformer framework that couples transformer-based semantic encoding with graph-based message passing. We represent the full corpus (train/validation/test) as nodes in a title-similarity graph built from bert-base-uncased title embeddings; to obtain a more informative similarity geometry, we apply PCA before computing similarities (thresholding can be used as an optional sparsification step). During training, transformer [CLS] embeddings are projected to a compact space and injected into a two-layer **Graph Convolutional Network (GCN)**, enabling information propagation from supervised nodes to the unlabeled portion of the corpus. Supervision is applied only on a feasible subset of labeled positive pairs and sampled negatives via a cosine embedding objective, while the remaining unlabeled reports contribute indirectly through graph propagation. For evaluation, model parameters are frozen and pairwise predictions are obtained by thresholding cosine similarity, where the decision threshold is selected on the validation split.

*Authors are listed in alphabetical order by surname.

Unpublished working draft. Not for distribution.
Permissions to copy or distribute in whole or in part in electronic or hard copy formats must be obtained from the author(s) or the copyright holder(s). Requests for permission to copy or redistribute should be addressed to permissions@acm.org.

CS 588, Fall 2025-26, Bilkent University

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXXXXXXXXX>

2026-01-03 10:30. Page 1 of 1-13.

Results. Experiments on Eclipse Platform and Mozilla Thunderbird show that graph propagation enables the model to benefit from unlabeled reports under label-scarce conditions, yielding stable training dynamics and competitive classification performance relative to transformer-only baselines. Source code and data are publicly available at the project repository.¹

Conclusions. The proposed graph-enhanced training strategy offers a principled way to exploit unlabeled bug reports without requiring explicit pairwise supervision for transformers. By confining graph-based reasoning to the transductive training phase, the framework balances label efficiency, scalability, and practical deployment constraints. These results indicate that graph-structured semi-supervised learning can effectively complement transformer-based encoders for duplicate bug report detection in realistic, label-scarce settings.

KEYWORDS

Bug Report, Bug Duplicate Detection, Transformers, BERT, LLMs, Graph Neural Networks, Semi-Supervised Learning

ACM Reference Format:

Kerem Bayramoğlu, Hüseyin Karaca, Emirhan Koç, Hasan Erkin Ünlü, and Rabia Ela Ünlü. 2018. Graph Enhanced Transformer for Semi-Supervised Duplicate Bug Report Detection. In *Proceedings of Data Science For Software Engineering (CS 588)*. ACM, New York, NY, USA, 13 pages. <https://doi.org/XXXXXXXXXXXXXX>

1 INTRODUCTION

In large software developments, bug-tracking systems are used as a tool for handling software bug reports, as well as handling software feature requests. In the course of time, as software **bug repositories** databases continue to expand, it has been noticed that there exist overlapping software bug reports from diverse users. This results in the enlargement of the size of the software bug-reporting databases, as well as consuming the time of software developers in examining the same software bugs, as these overlapping software bug reports are nothing but duplicates. Studies show that a large number of

¹<https://github.com/huseyin-karaca/graph-enhanced-dbd.git>

117 duplicate bug reports exist in repositories; for example, 20% of
 118 reports in Eclipse and 30% in Firefox were marked as duplicate [2].

119 Automating duplicate bug report detection has therefore become
 120 an essential research direction in software engineering. Classical in-
 121 formation retrieval (IR) techniques such as Term Frequency–Inverse
 122 Document Frequency (TF-IDF), Best Matching 25 with field weight-
 123 ing (BM25F), and custom retrieval functions (e.g., REP) have been
 124 widely used to match new bug reports with existing ones [32].

125 Although these algorithms are computationally efficient, they
 126 remain limited in addressing the vocabulary mismatch problem,
 127 where semantically similar bug reports are expressed using different
 128 lexical choices [8]. More recent approaches based on deep learning,
 129 particularly transformer-based models derived from BERT, have
 130 demonstrated improved effectiveness in capturing semantic similar-
 131 ity by leveraging contextual representations of bug report titles and
 132 descriptions [5, 20, 27]. Despite these advances, transformer-based
 133 methods rely heavily on supervised learning with labeled duplicate
 134 reports, which introduces a substantial dependency on annotated
 135 data. In practice, acquiring such labels is costly and labor-intensive,
 136 and real-world bug repositories typically contain only a limited
 137 number of confirmed duplicate annotations. To mitigate the re-
 138 liance on limited annotated data, several studies have proposed
 139 augmenting supervision by exploiting unlabeled bug reports, for in-
 140 stance through pairing strategies combined with negative sampling
 141 [22, 26]. While such approaches partially alleviate label scarcity,
 142 they remain limited in their ability to fully leverage the entire pool
 143 of available unlabeled data, and thus do not provide a comprehen-
 144 sive solution for large-scale, label-sparse bug repositories.

145 To address these limitations, we propose a semi-supervised Graph-
 146 Enhanced Transformer framework for duplicate bug report detec-
 147 tion. Our approach leverages the semantic power of transformer
 148 encoders, along with the relational reasoning power of graph neu-
 149 ral networks (GNNs). Unlike transformer-based models, which are
 150 inherently limited in their ability to fully exploit unlabeled data, our
 151 framework explicitly incorporates both labeled and unlabeled bug
 152 reports as nodes within a unified graph structure. By leveraging the
 153 message-passing mechanism of GNNs, information is propagated
 154 between labeled and unlabeled nodes, enabling implicit knowledge
 155 transfer across the entire report collection. Although the training
 156 objective and final predictions are defined solely over labeled data,
 157 the inclusion of unlabeled nodes allows the model to benefit from
 158 their structural and semantic context, preventing their complete
 159 exclusion from the learning process.

160 This study is guided by the following research questions:

- 161 • **RQ1:** Can unlabeled bug reports be effectively leveraged
 162 through graph-based representations for duplicate bug re-
 163 port detection?
- 164 • **RQ2:** Can a scalable architecture be designed to support du-
 165 plicate bug report detection in large-scale bug repositories?
- 166 • **RQ3:** Does the proposed graph-enhanced, semi-supervised
 167 framework achieve competitive performance compared to
 168 strong transformer-based baselines?

169 This study offers the following contributions:

- 170 • We explicitly model both **labeled and unlabeled bug re-**
 171 **ports as nodes in a unified graph**, enabling effective ex-
 172 **ploitation of unlabeled data through graph-based message**

175 passing, while defining supervision and prediction solely on
 176 labeled samples.

- 177 • We show that unlabeled bug reports contribute indirectly to
 178 learning by providing **structural and semantic context**,
 179 facilitating implicit knowledge transfer across the entire
 180 report collection.
- 181 • To the best of our knowledge, this work represents **an early**
 182 **exploration of GNN-based approaches for duplicate**
 183 **bug report detection**, revealing the potential of graph-
 184 enhanced learning in this domain.
- 185 • Although the proposed framework adopts a **transductive**
 186 **training and inference setting**, it is **not inherently re-**
 187 **stricted to this regime** and can be naturally extended to
 188 **inductive inference** with appropriate architectural modifi-
 189 **cations**.

2 PROBLEM DESCRIPTION AND MOTIVATION

2.1 Formal Problem Definition

The core challenge in Duplicate Bug Report Detection (DBRD) is to automatically identify whether two bug reports describe the same underlying software defect. Formally, given a bug repository $\mathcal{R} = \{r_1, r_2, \dots, r_n\}$ where each report r_i contains textual fields (title, description) and metadata (timestamp, reporter, component), the task is to learn a function $f : \mathcal{R} \times \mathcal{R} \rightarrow \{0, 1\}$ that predicts whether a pair of reports (r_i, r_j) are duplicates. This is inherently difficult due to lexical variations, incomplete descriptions, domain-specific terminology, and reporting style differences across users.

In realistic software development environments, the availability of labeled data is severely limited. While a repository may contain tens of thousands of bug reports, only a small fraction (typically less than 5%) have confirmed duplicate relationships manually identified by developers or triagers. The remaining reports remain unlabeled, representing a vast untapped resource that conventional supervised learning approaches cannot effectively utilize.

2.2 Limitations of Existing LLM-Based Approaches

Transformer-based language models such as BERT, RoBERTa, and their domain-adapted variants have demonstrated strong performance in semantic text matching tasks. However, these models face fundamental limitations in the DBRD setting. Consider a concrete scenario: a repository contains 20,000 bug reports with only 500 confirmed duplicate pairs (1,000 labeled reports). A pairwise BERT-based classifier would be trained on these 500 positive pairs plus an equal number of sampled negative pairs, effectively utilizing only 1,000 reports while completely disregarding the remaining 19,000 unlabeled reports.

The quadratic growth of potential pairs exacerbates this problem. For n bug reports, there exist $\binom{n}{2} = \frac{n(n-1)}{2}$ possible pairs. For 20,000 reports, this yields approximately 200 million pairs, making exhaustive pairwise training computationally infeasible. Existing approaches employ negative sampling strategies, but these methods lack a principled mechanism to propagate semantic information across the unlabeled corpus. Moreover, they are limited by the

233 explicit supervision paradigm: without a label for a pair (r_i, r_j) , the
 234 model receives no training signal involving those reports together.
 235
 236
 237

2.3 Illustrative Example: LLM Failure Under Data Scarcity

240 To illustrate the practical limitation, consider two bug reports from
 241 an Eclipse repository: *Report A (ID 12345)*: "NullPointerException
 242 in editor when opening XML file with invalid schema reference"
 243 *Report B (ID 67890)*: "NPE thrown during XML editor initialization
 244 with broken schema link." These reports clearly describe the same
 245 bug using different terminology ("NullPointerException" vs. "NPE",
 246 "opening" vs. "initialization", "invalid reference" vs. "broken link").
 247 A well-trained transformer model with sufficient labeled examples
 248 from the XML editor domain would likely identify this pair as duplicates
 249 through semantic similarity. However, suppose the training
 250 set contains only 100 labeled duplicate pairs, none of which involve
 251 the XML editor component or schema-related issues. The trans-
 252 former model, trained exclusively on these 100 pairs, would lack
 253 the domain-specific context to recognize the semantic equivalence
 254 between Reports A and B. Meanwhile, the repository contains 500
 255 other unlabeled reports related to XML editing, including variations
 256 of similar terminology and error patterns. A purely supervised LLM-
 257 based approach cannot leverage these 500 reports during training,
 258 as they lack explicit duplicate labels.

2.4 Motivation for Graph-Based Semi-Supervised Learning

261 Graph neural networks provide a natural solution to overcome these
 262 limitations. Unlike pairwise supervised learning, GNNs operate
 263 on relational neighborhoods and propagate information through
 264 message passing across connected nodes without requiring explicit
 265 labels for every connection. This enables us to construct a unified
 266 graph representation where all available bug reports—both labeled
 267 and unlabeled—participate in the learning process.

268 In our framework, edges are formed based on label-independent
 269 semantic similarity (e.g., cosine similarity of title embeddings), en-
 270 abling efficient graph construction without additional manual an-
 271 notations. The transformer component operates on a restricted but
 272 feasible subset of report pairs during training: positive pairs from
 273 known duplicates and negative pairs sampled from different groups.
 274 The representations learned from these labeled pairs are embedded
 275 as graph node features, and the GNN propagates this information
 276 across the entire graph through neighborhood aggregation.

277 Returning to the XML editor example: even though Reports A
 278 and B may not appear in the training set, if they are connected
 279 through intermediate nodes (other XML-related reports) in the
 280 graph, the GNN can propagate relevant semantic information to
 281 them. Although direct supervision is applied only to nodes in la-
 282 beled pairs, the graph structure allows information flow to unpaired
 283 and unlabeled nodes, enabling the model to benefit from the full cor-
 284 pus. This design decouples pairwise supervision from global data
 285 utilization, making it suitable for realistic, label-scarce duplicate
 286 bug report detection scenarios.

3 RELATED WORK

291 The problem of duplicate bug report detection has been a central
 292 challenge in software engineering research for many years. Early
 293 approaches primarily relied on classical Information Retrieval (IR)
 294 methods, while more recent techniques have leveraged advances
 295 in machine learning and deep learning to improve performance. In
 296 this section, we review the evolution of methodologies for DBRD,
 297 highlighting key contributions and their limitations.

3.1 Information Retrieval Models

298 Early automated approaches framed DBRD as a classical Informa-
 299 tion Retrieval (IR) problem [33], typically using the Vector Space
 300 Model (VSM) [23]. In this model, each bug report is treated as a
 301 document and is transformed into a high-dimensional vector. The
 302 components of this vector are weighted based on the terms present
 303 in the document.

304 The most common scheme is Term Frequency-Inverse Document
 305 Frequency (TF-IDF) [29, 31]. This method assigns high weights to
 306 terms that are frequent in a specific document but rare across the
 307 entire corpus. Once vectorized, report similarity is calculated using
 308 Cosine Similarity [23].

309 The fundamental limitation of this paradigm is the “vocabulary
 310 mismatch problem” [9]. First identified by Furnas et al. [9], this
 311 refers to the low probability that two people will use the same
 312 terms for the same concept.

3.2 Machine Learning Approaches

313 In response, researchers applied supervised machine learning, shift-
 314 ing the problem from retrieval to classification (a binary duplicate/non-
 315 duplicate decision) [12, 33]. Early work, such as Jalbert and Weimer
 316 (2008) [12] and Sun et al. (2010) [33], applied classifiers like Support
 317 Vector Machines (SVMs) to feature pairs.

318 However, training a classifier requires generating $O(N^2)$ poten-
 319 tial pairs, where N denotes the total number of bug reports in the
 320 repository, which is computationally intractable [33]. This “low
 321 efficiency” of pair-wise classification led to a critical development:
 322 using machine learning to improve the IR model itself.

323 This led to REP, a retrieval function proposed by Sun et al. [32]
 324 that became a dominant baseline. REP extends the Okapi BM25F
 325 formula to create a custom similarity function. Its key innovation
 326 is combining textual similarity with structured metadata (product,
 327 component, version) [32]. It then uses supervised learning to learn
 328 the optimal weights for each field, tuning its function for a specific
 329 bug repository [32]. The success of REP demonstrated that a spe-
 330 cialized, feature-aware IR model could outperform general-purpose
 331 ML models in both accuracy and efficiency [32, 33], and it remained
 332 a difficult baseline to beat for years [32].

3.3 Deep Learning for Semantic Representation

333 The deep learning (DL) revolution promised to finally solve the
 334 vocabulary mismatch problem by learning true semantic meaning
 335 [3]. Initial attempts with Convolutional Neural Networks (CNNs)
 336 and Recurrent Neural Networks (RNNs), such as the Dual-Channel
 337 CNN (DCCNN) by He et al. (2020) [11], began learning vector
 338 embeddings from report text. However, these early DL models

339
 340
 341
 342
 343
 344
 345
 346
 347
 348

often struggled to outperform the highly-optimized REP baseline [32].

A significant shift occurred with the Transformer model, specifically BERT [6]. BERT's pre-training allows it to capture deep, contextual understanding of language [6]. This base model was quickly followed by optimized variants such as RoBERTa [21], which improved performance by refining the pre-training process, and ALBERT [19], which focused on parameter reduction. However, applying these cross-encoder models (BERT, RoBERTa, ALBERT) directly to DBRD exposed a critical flaw. As a cross-encoder, BERT requires both reports to be fed into the network simultaneously [28]. Finding the best match for a new report in a large database would require 10,000s of computations, a process estimated to take around 65 hours [28], rendering it unusable.

This was solved by Reimers and Gurevych (2019) with Sentence-BERT (SBERT) [28]. SBERT adapts BERT using a siamese network structure [28]. Two identical BERT models process each bug report independently, producing a single "sentence embedding" for the pair [28]. Because embeddings are generated independently, they can be pre-computed. A new report can be embedded once, and its similarity to all others found almost instantaneously via cosine similarity search [28]. This reduces the 65-hour search task to approximately 5 seconds [28]. SBERT provides a powerful, semantic-native replacement for TF-IDF, solidifying the dominance of the two-stage "retrieval-rerank" pipeline. A recent comparative study by Meng et al. (2024) empirically evaluated many of these architectures, confirming the performance gains of transformer-based models like BERT, ALBERT, and RoBERTa over earlier neural architectures like Bi-LSTM and DC-CNN [24].

3.4 Hybrid Systems and Emerging Approaches

Current state-of-the-art systems are often hybrid approaches combining lexical precision with semantic understanding [25]. The DBTM (Duplicate Bug report Topic Model) approach, for example, combines the IR model BM25F with topic-based features [25], allowing it to match reports based on higher-level "technical bugs" [25].

More recently, this philosophy has extended to Large Language Models (LLMs). Cupid [34] achieves state-of-the-art results by combining an LLM (ChatGPT) with the classic IR model REP [32, 34]. The LLM is used in a zero-shot setting as a semantic pre-processor to "get essential information" from the raw bug report [34]. This "cleaned" information is then fed to the domain-aware REP model for the final similarity calculation. This hybrid approach was shown to improve recall over previous baselines by a significant margin [34].

While the aforementioned models have progressively advanced semantic understanding, they remain almost exclusively *supervised*, requiring large, costly datasets of labeled duplicate pairs. As noted in recent studies, transformer models like BERT and its variants (ALBERT, RoBERTa) excel when data is plentiful, but their performance is limited in the more realistic, label-scarce environments common to bug repositories [24]. Furthermore, as GNNs emerge as a new frontier, existing applications are often *transductive*, meaning they cannot perform inference on new, unseen bug reports without retraining [10].

Our proposed framework addresses these two specific, critical gaps: (1) label scarcity and (2) inference scalability. We operate in a *semi-supervised* setting, using a GNN to leverage the vast majority of *unlabeled* reports during training. Table 1 outlines this positioning relative to closely related works.

Stepping back, we view DBRD as both a semantic and a relational problem. In the next section, we model the bug repository as a graph and use an inductive GNN (e.g., GraphSAGE) to produce embeddings for unseen reports without retraining [10]. To deal with limited supervision and sharpen the decision boundary, we add two simple training aids: generative augmentation to create within-class variants [1] and hard-negative mining to focus the model on near-miss non-duplicates. We detail the architecture and training strategy in the next section.

4 PROPOSED METHOD

We introduce a semi-supervised GNN augmented transformer-encoder framework for identification of duplicate bug reports, where the model is trained primarily in positive-pair supervision with a regularization strategy, i.e., the negative pair sampling . Overall, the pipeline consists of three main stages: (i) Data Processing, where duplicate labels are transformed into positive training pairs and all remaining pairs are treated as candidate negatives; (ii) Embedding Construction, Training as well as Validation, where a transformer encoder learns similarity-preserving representations using a binary similarity objective; and (iii) Inference, where the model is tested on untrained positive and negative pairs. The overall schematic encompassing data processing, embedding construction, information transfer to the GNN and training is illustrated in Fig. 1. The description and formal definition of each stage is provided as follows:

Pair Construction. Let $\mathcal{P} = \{p_1, p_2, \dots, p_N\}$ denote the set of all bug reports, where each report p_i is associated with a duplicate-group label $g_i \in \{1, \dots, L\}$.

Positive pairs. For each report index $i \in \{1, 2, \dots, N\}$, we define the set of positive (duplicate) pairs as

$$\mathcal{P}_i^{(+)} = \{(i, j) \mid g_i = g_j, j \neq i\}.$$

These pairs enforce similarity among reports belonging to the same duplicate group.

Anchor-based negative pairs. For each report index $i \in \{1, 2, \dots, N\}$, we define the anchor-based negative pair set as

$$\mathcal{N}_i^{(a)} = \{(i, j) \mid g_i \neq g_j\}.$$

Since $|\mathcal{N}_i^{(a)}|$ is typically much larger than $|\mathcal{P}_i^{(+)}|$, we randomly sample a subset

$$\tilde{\mathcal{N}}_i^{(a)} \subset \mathcal{N}_i^{(a)}, \quad |\tilde{\mathcal{N}}_i^{(a)}| = \alpha |\mathcal{N}_i^{(a)}|,$$

where α denotes the fixed subsampling ratio.

Cross-group random negative pairs. To further diversify negative supervision, we additionally construct anchor-free negative pairs by randomly sampling report pairs across different duplicate groups:

$$\mathcal{N}^{(r)} = \{(i, j) \mid i \neq j, g_i \neq g_j\}.$$

From this set, a random subset $\tilde{\mathcal{N}}^{(r)} \subset \mathcal{N}^{(r)}$ is selected.

Table 1: Comparison of Closely Related DBRD Methodologies

Methodology	Core Technology	Semantic Capacity	Uses Unlabeled Data?	Inference Scalability
Bi-LSTM [4]	Recurrent Neural Network (RNN)	Sequential	No (Supervised)	High (Bi-encoder)
DC-CNN [11]	Convolutional Neural Network (CNN)	Local Patterns	No (Supervised)	High (Bi-encoder)
BERT/ALBERT/RoBERTa [6, 19, 21]	Transformer (Cross-Encoder)	Deep Contextual	No (Supervised)	Low (Pair-wise)
SBERT [28]	Transformer (Siamese Bi-Encoder)	Deep Contextual	No (Supervised)	High (Bi-encoder)
Our Method	Transformer + GNN (Hybrid)	Deep Contextual & Relational	Yes (Semi-supervised)	High (Bi-encoder)

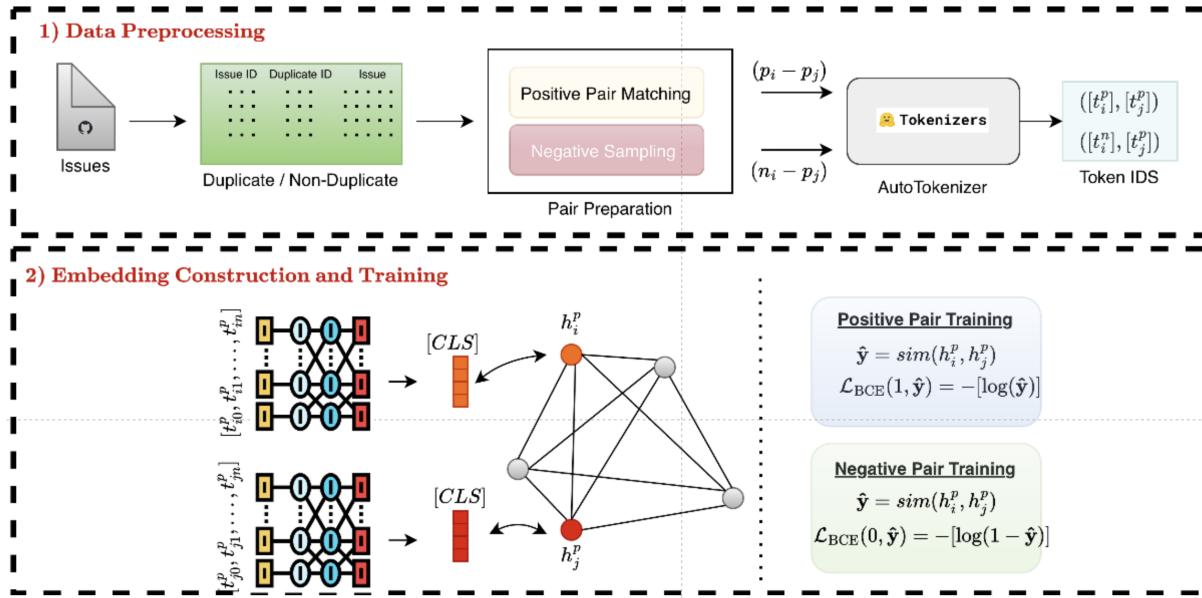


Figure 1: Overall pipeline of the proposed method. Upper: The bugs are documented with bug numbers, titles, and descriptions. The positive and negative pairs are collected and stacked. Then, the token IDs are listed to be fed into BERT. Bottom: The graph is constructed, then BERT is employed for feature extractor and these embeddings transferred to the GNN component. Finally, the model is trained based on similarity and dissimilarities between pairs.

Final training pair sets. The final positive and negative pair sets used during training are

$$\mathcal{P}^{(+)} = \bigcup_{i=1}^N \mathcal{P}_i^{(+)}, \quad N = \bigcup_{i=1}^N \tilde{\mathcal{N}}_i^{(a)} \cup \tilde{\mathcal{N}}^{(r)}.$$

Following pair construction, each bug report is converted into a token-level representation using a HuggingFace tokenizer. Let x_i denote the textual content of report p_i , obtained by concatenating its title and description.

For each positive pair $(i, j) \in \mathcal{P}_i^{(+)}$ and each sampled negative pair $(i, j) \in \tilde{\mathcal{N}}_i^{(a)} \cup \tilde{\mathcal{N}}^{(r)}$, we obtain the corresponding token-ID sequences as

$$t_i = \text{Tokenizer}(x_i), \quad t_j = \text{Tokenizer}(x_j).$$

The same tokenization process is applied to both positive and negative pairs, ensuring a unified input format for the transformer encoder. Each token-ID sequence is then used as direct input to the transformer-based embedding module, which is described in the following sections.

4.1 Embedding and Graph Construction

In the embedding construction stage, our methodology consists of three fundamental components: *Graph Construction*, *Embedding Construction*, and *End-to-End Training*.

4.1.1 Graph Construction. GNNs provide a flexible representation learning framework capable of operating on both labeled and unlabeled nodes. A GNN defines a message-passing operator over a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where information is iteratively exchanged among neighboring nodes to produce context-aware node embeddings [13, 14]. Since the update functions in a GNN do not rely on node labels, the model naturally supports semi-supervised settings in which only a subset of nodes is labeled while the remaining nodes remain unlabeled [16]. This property is particularly well aligned with our setting, where duplicate-group information is available for only a small fraction of reports.

We construct the graph using the entire dataset, such that every bug in the corpus, including training, validation, and test samples, is represented as a node. Each node corresponds to a unique bug report, and node indices are explicitly preserved to ensure a consistent mapping between reports and their graph representations throughout training and inference. Although GNNs operate in a

581 permutation-invariant manner, maintaining this index correspondence is essential for correctly associating learned representations
 582 with their respective reports. Edges in the graph encode pairwise
 583 semantic relations between bugs; while each edge can be interpreted
 584 as a potential report pair, direct connectivity between all
 585 relevant nodes is not required, as information can propagate across
 586 the graph through multi-hop message passing.
 587

In our graph construction, each bug is assigned an initial node
 588 feature vector corresponding to a one-hot encoding of its node
 589 index. Formally, for a graph with N nodes, the initial embedding of
 590 node i is given by
 591

$$x_i^{(0)} = e_i \in \mathbb{R}^N,$$

592 where e_i denotes the i -th standard basis vector.
 593

The edge set is constructed using semantic similarity computed
 594 from bug titles. Let t_i denote the tokenized title of bug i . Each title
 595 is first encoded using a pretrained bert-base-uncased model,
 596 producing a contextual embedding in \mathbb{R}^{768} . Based on empirical
 597 observations, similarity computations performed directly in this
 598 high-dimensional space tend to be overly smooth and less infor-
 599 mative for graph construction. To mitigate this effect, we apply
 600 Principal Component Analysis (PCA) to the title embeddings and
 601 retain the top $d = 10$ principal components, resulting in reduced rep-
 602 resentations $\tilde{t}_i \in \mathbb{R}^{10}$. Similarity is then computed in this reduced
 603 space, yielding more discriminative and semantically meaningful
 604 relationships by alleviating redundancy and noise in the original
 605 embeddings.
 606

Edges are added between nodes based on the resulting similarity
 607 structure, without enforcing a hard threshold. Instead, similarity
 608 values are used directly to define graph connectivity, enabling
 609 flexible and dense relational modeling across the dataset. While
 610 threshold-based edge selection constitutes a viable alternative for
 611 graph sparsification, our formulation does not rely on an explicit
 612 similarity cutoff.
 613

614 **4.1.2 Embedding Construction.** For each token-ID sequence ob-
 615 tained in the data preprocessing stage, the embedding construction
 616 module feeds the sequence into a pre-trained transformer encoder.
 617 Each element of a report pair is processed independently through
 618 the encoder, yielding two contextualized representations. In prac-
 619 tice, this corresponds to two parallel forward passes through a pair
 620 of parameter-tied transformer encoders, ensuring that both inputs
 621 are mapped into a shared representation space. These operations
 622 are applied uniformly to both positive and negative pair elements.
 623

Formally, given a token-ID sequence
 624

$$t^{(i)} = [t_0^{(i)}, \dots, t_{L_i}^{(i)}],$$

the pretrained transformer encoder is denoted by
 625

$$f_W : \mathbb{N}^{L_i} \rightarrow \mathbb{R}^{L_i \times d},$$

630 where W represents the shared parameters of the weight-tied en-
 631 coders and d denotes the hidden dimension of the transformer. The
 632 resulting contextual embedding matrix is
 633

$$H_i = f_W(t^{(i)}) \in \mathbb{R}^{L_i \times d}.$$

The report-level embedding is extracted from the special [CLS]
 635 token position,
 636

$$h_i^{(\text{CLS})} = H_i[0] \in \mathbb{R}^d,$$

637 which serves as a pooled summary representation of the entire
 638 input sequence [16, 17].
 639

To reduce computational and memory overhead, particularly
 640 under small-scale computing constraints, the high-dimensional
 641 [CLS] embeddings are further projected into a lower-dimensional
 642 space using a learnable linear transformation:
 643

$$z_i = W_p h_i^{(\text{CLS})} + b_p, \quad W_p \in \mathbb{R}^{D \times d},$$

644 yielding compact embeddings $z_i \in \mathbb{R}^D$. In our experiments, D is
 645 treated as a tunable hyperparameter, with $D = 128$ used as a stable
 646 operating point unless otherwise stated.
 647

648 For each supervised pair $(i, j) \in \mathcal{S}$, the final embeddings are
 649 obtained via two parallel forward passes followed by projection:
 650

$$z_i = W_p f_W(t^{(i)})[0], \quad z_j = W_p f_W(t^{(j)})[0].$$

651 Parameter sharing across both the transformer encoder and the
 652 projection layer ensures that all reports are embedded into a con-
 653 sistent low-dimensional representation space suitable for efficient
 654 similarity-based training and inference.
 655

656 **4.1.3 Training Pipeline.** In the training pipeline, we combine the
 657 outputs of the transformer-based embedding module and the graph
 658 neural network. The token-ID sequences obtained from the report
 659 descriptions are first fed into the pretrained transformer encoder,
 660 producing the corresponding [CLS] embeddings. For a report pair
 661 (i, j) , let h_i and h_j denote the transformer-derived [CLS] embed-
 662 dings. These embeddings are then injected into the GNN by assign-
 663 ing them to their corresponding nodes in the constructed graph,
 664 ensuring consistency with the fixed node ordering. Importantly,
 665 supervision is applied only to nodes whose indices belong to the
 666 training split, while validation and test nodes are included in the
 667 graph to enable information propagation via message passing.
 668

669 **Fusion of Transformer and GNN representations.** Let $z_i^{(\text{tr})} \in \mathbb{R}^D$
 670 denote the projected transformer embedding of report i , and let
 671 $z_i^{(\text{gnn})} \in \mathbb{R}^D$ denote the corresponding GNN output embedding
 672 after message passing. The final embedding used for optimization
 673 is obtained via a weighted fusion:
 674

$$z_i = \lambda z_i^{(\text{tr})} + (1 - \lambda) z_i^{(\text{gnn})}, \quad \lambda \in [0, 1].$$

675 This formulation balances semantic information captured by the
 676 transformer with relational information captured by the GNN. As
 677 an alternative design choice, the two embeddings may also be com-
 678 bined via concatenation, i.e., $z_i = [z_i^{(\text{tr})}; z_i^{(\text{gnn})}]$, followed by a
 679 projection layer. Unless otherwise stated, the weighted-sum fusion
 680 is used throughout this work.
 681

682 **Cosine embedding objective.** For each supervised report pair (i, j) ,
 683 we compute the cosine similarity
 684

$$s_{ij} = \cos(z_i, z_j).$$

685 We optimize a cosine embedding loss, where pairwise labels are
 686 mapped to $\{-1, +1\}$. Specifically, positive (duplicate) pairs are as-
 687 signed $y_{ij} = +1$, while sampled negative pairs are assigned $y_{ij} = -1$.
 688 The loss for a pair (i, j) is defined as
 689

$$\mathcal{L}_{ij} = \begin{cases} 1 - s_{ij}, & y_{ij} = +1, \\ \max(0, s_{ij} - m), & y_{ij} = -1, \end{cases}$$

where $m \in [0, 1]$ denotes a margin hyperparameter. The overall training objective is given by

$$\mathcal{L} = \sum_{(i,j) \in \mathcal{P}^{(+)}} \mathcal{L}_{ij} + \sum_{(i,j) \in \mathcal{N}} \mathcal{L}_{ij}.$$

Minimizing \mathcal{L} jointly updates the parameters of the transformer encoder, the projection module, and the GNN, enabling similarity-preserving representations informed by both textual semantics and graph-structured relational information.

Validation-based decision rule. During validation, similarity scores are computed only for pairs involving nodes belonging to the validation split. A pair (i, j) is predicted as duplicate if its similarity score exceeds a threshold θ :

$$\hat{y}_{ij} = \mathbb{I}[s_{ij} \geq \theta].$$

A default threshold of $\theta = 0.5$ is used as an initial operating point during training. To determine an appropriate decision threshold, we perform a validation-time grid search over $\theta \in [0, 1]$ with a step size of 0.05, and select the value that yields the best validation performance. The selected threshold is then fixed for subsequent evaluation.

4.2 Inference Procedure

Let \hat{f}_W and \hat{g}_Θ denote the transformer encoder (including projection) and the GNN with parameters fixed after training. During inference, all model parameters are frozen and no further updates are performed.

Pairwise Inference on Test Nodes. For test nodes indexed by $\mathcal{I}_{\text{test}}$, we follow a procedure analogous to validation. For each test pair (i, j) with $i, j \in \mathcal{I}_{\text{test}}$, embeddings are computed as

$$z_i = \hat{f}_W(t^{(i)}), \quad z_j = \hat{f}_W(t^{(j)}),$$

optionally refined via graph propagation using \hat{g}_Θ . The cosine similarity

$$s_{ij} = \cos(z_i, z_j)$$

is then compared against the decision threshold θ , yielding the prediction

$$\hat{y}_{ij} = \mathbb{I}[s_{ij} \geq \theta].$$

This pairwise evaluation protocol is adopted to ensure consistency with existing benchmarks and prior work.

5 EXPERIMENTS AND RESULTS

5.1 Dataset

Our experimental evaluation employs two prevalently utilized benchmark datasets for duplicate bug report detection. Eclipse Platform, and Mozilla Thunderbird [18]. These include both original and duplicate reports from large, long duration software projects. It can be referred to Table 2 for details.

5.2 Experimental Setup

5.2.1 Evaluation Metrics. Let $\mathcal{T} \subseteq \mathcal{I}_{\text{test}} \times \mathcal{I}_{\text{test}}$ denote the set of evaluated test pairs, and let $y_{ij} \in \{0, 1\}$ and $\hat{y}_{ij} \in \{0, 1\}$ denote the ground-truth and predicted duplicate labels for a pair $(i, j) \in \mathcal{T}$,

Table 2: Statistics of the datasets used for duplicate bug report detection.

Metric	Eclipse	Thunderbird
Total	68,124	26,040
Non-Dup.	50,606	15,994
Dup.	17,512	10,046
% Dup.	25.7	38.6
Clusters	6,282	2,772
Pairs	87,224	53,156

respectively. We define the sets of true positives, false positives, true negatives, and false negatives as

$$\begin{aligned} \text{TP} &= \{(i, j) \in \mathcal{T} \mid y_{ij} = 1 \wedge \hat{y}_{ij} = 1\}, \\ \text{FP} &= \{(i, j) \in \mathcal{T} \mid y_{ij} = 0 \wedge \hat{y}_{ij} = 1\}, \\ \text{TN} &= \{(i, j) \in \mathcal{T} \mid y_{ij} = 0 \wedge \hat{y}_{ij} = 0\}, \\ \text{FN} &= \{(i, j) \in \mathcal{T} \mid y_{ij} = 1 \wedge \hat{y}_{ij} = 0\}. \end{aligned}$$

Using these quantities, accuracy is computed as

$$\text{Accuracy} = \frac{|\text{TP}| + |\text{TN}|}{|\text{TP}| + |\text{TN}| + |\text{FP}| + |\text{FN}|}.$$

Precision and recall are defined as

$$\text{Precision} = \frac{|\text{TP}|}{|\text{TP}| + |\text{FP}|}, \quad \text{Recall} = \frac{|\text{TP}|}{|\text{TP}| + |\text{FN}|},$$

and the F1-score is given by

$$\text{F1} = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}.$$

5.2.2 Implementation Details. Our implementation is built using PyTorch and the HuggingFace Transformers library. For the transformer encoder, we experiment with four pre-trained models: BERT [5], RoBERTa [21], DistilBERT [30], and CodeBERT [7]. The graph neural network component employs a two-layer Graph Convolutional Network (GCN) architecture [15].

For training, we use the AdamW optimizer with a learning rate of 2×10^{-5} and a batch size of 64. Models are trained for up to 5 epochs with early stopping based on validation performance; specifically, training is stopped if validation loss does not improve for 3 consecutive epochs (patience=3). Negative sampling is performed using two complementary strategies. First, anchor-based negative pairs are sampled for each report in proportions comparable to the number of positive pairs. Second, to increase diversity, an additional set of 5,000 negative pairs is randomly sampled across different duplicate groups. Importantly, the vast majority of unlabeled reports are not included in transformer-based pairwise training; instead, they are incorporated exclusively as nodes in the graph, enabling them to influence learning through GNN-based message passing without being explicitly paired.

Graph construction is performed using semantic similarity computed from bug titles encoded with a pretrained bert-base-uncased model. Similarity is computed directly on these title embeddings after PCA-based dimensionality reduction, and no explicit threshold is applied for edge selection. This design choice allows dense

and flexible connectivity in the graph while avoiding sensitivity to threshold tuning.

All experiments were conducted on a single NVIDIA high performance A100 GPU with 80 GB VRAM. All source code and the full reproducibility package, including all necessary data files and scripts, can be found in the [huseyin-karaca/graph-enhanced-dbd GitHub repository](https://github.com/huseyin-karaca/graph-enhanced-dbd).

5.3 Results

We organize our experimental results around the three research questions presented in Section 1.

5.3.1 Leveraging Unlabeled Data Through Graph Structure (RQ1). To investigate whether unlabeled bug reports can be effectively leveraged through graph-based representations, we analyze the similarity distributions in the learned embedding space. Table 3 presents representative examples from the validation set, showing the cosine similarity scores for both duplicate and non-duplicate pairs.

Table 3: Representative similarity scores on validation set pairs. Duplicate pairs consistently achieve high similarity, while non-duplicate pairs show low or negative similarity.

Duplicates (Label = 1)			Non-Duplicates (Label = 0)		
Issue 1	Issue 2	Similarity	Issue 1	Issue 2	Similarity
8470	35054	0.923	158896	47204	-0.033
70983	83204	0.993	124062	21003	-0.107
62405	104926	0.979	73654	28863	0.138

The results demonstrate clear separation in the learned embedding space. Duplicate pairs consistently achieve high cosine similarity scores, while non-duplicate pairs exhibit significantly lower similarities. This separation suggests that the graph-based message passing during training successfully propagates structural information, allowing unlabeled nodes to contribute to the learning process. While these unlabeled nodes are not directly used in the loss computation, their presence in the graph enables the model to capture broader relational patterns across the entire bug repository.

Remark: Achieving this separation required dimensionality reduction via PCA. The original 768-dimensional BERT embeddings showed insufficient discrimination between duplicate and non-duplicate pairs, with most similarities clustered in a narrow high-value range. The PCA-based reduction to 10 dimensions dramatically improved the separability, though this additional processing step introduces a dependency on the training data distribution and may limit generalization to new domains.

Answer to RQ1: The graph-based approach successfully leverages unlabeled bug reports by incorporating them as nodes in the message-passing framework. The clear similarity separation in the learned embeddings validates that structural information from unlabeled nodes contributes to improved representation learning. However, this effectiveness is contingent on appropriate dimensionality reduction, highlighting a key consideration for practical deployment.

5.3.2 Scalability of the Architecture (RQ2). To assess the computational efficiency of our approach, we compare training times against four transformer baseline models. Table 4 presents the results on the Eclipse and Thunderbird datasets.

Table 4: Training time comparison on Eclipse and Thunderbird datasets. Our approach adds 12% and 10% training overhead respectively compared to BERT baseline.

Model	Eclipse (min/epoch)		Thunderbird (min/epoch)	
	Time	Rel.	Time	Rel.
BERT [5]	16.2	1.0×	10.1	1.0×
RoBERTa [21]	16.3	1.01×	10.2	1.01×
DistilBERT [30]	15.9	0.98×	9.8	0.97×
CodeBERT [7]	16.5	1.02×	10.3	1.02×
Ours	18.1	1.12×	11.1	1.10×

Our approach incurs a 12% training time overhead (18.1 minutes per epoch vs. 16.2 minutes for BERT) on Eclipse dataset due to the additional graph message-passing operations. On Mozilla Thunderbird dataset, our model takes 11.1 minutes per epoch, while the BERT model takes 10.1 minutes.

Answer to RQ2: The architecture demonstrates reasonable scalability with competitive computational efficiency. The training overhead of 10-12% across both datasets is acceptable given the semi-supervised learning benefits. The consistency of overhead across different dataset sizes (Eclipse with 68K samples and Thunderbird with 26K samples) indicates that our graph-enhanced approach scales proportionally without introducing disproportionate computational costs. The architecture is also available to successfully achieves the design goal of removing the GNN during inference, preventing the need for graph reconstruction with new bug reports.

5.3.3 Competitive Performance Evaluation (RQ3). Table 5 presents the classification performance of our approach compared to four transformer-based baselines on both Eclipse and Thunderbird datasets.

Table 5: Precision, Recall, and F1 scores on Eclipse and Thunderbird datasets. Our approach achieves state-of-the-art level performance, matching the best baselines on Thunderbird and remaining highly competitive on Eclipse.

Model	Eclipse			Thunderbird		
	P	R	F1	P	R	F1
BERT [5]	0.921	0.933	0.927	0.953	0.961	0.957
RoBERTa [21]	0.918	0.930	0.924	0.946	0.968	0.957
DistilBERT [30]	0.936	0.921	0.929	0.950	0.965	0.957
CodeBERT [7]	0.920	0.928	0.924	0.902	0.930	0.916
Ours	0.913	0.935	0.924	0.949	0.966	0.957

Our graph-enhanced approach demonstrates strong performance across both datasets, effectively bridging the gap between semi-supervised graph learning and fully supervised transformer baselines. On the Thunderbird dataset, our model achieves an F1 score of **0.957** (precision: 0.949, recall: 0.966), matching the top-performing


929 baselines (BERT, RoBERTa, and DistilBERT) exactly. This indicates
 930 that our method successfully captures the semantic nuances of du-
 931 plicate reports as effectively as heavy, fully-supervised pre-trained
 932 models.

933 On the Eclipse dataset, our model reaches an F1 score of **0.924**
 934 (precision: 0.913, recall: 0.935). This performance is on par with
 935 RoBERTa (0.924) and CodeBERT (0.924), and falls only marginally
 936 behind the highest-performing baseline, DistilBERT (0.929), by a
 937 negligible margin of 0.5%. Notably, our model exhibits the highest
 938 recall (0.935) on the Eclipse dataset among all comparison models,
 939 suggesting that the graph structural information aids significantly
 940 in retrieving relevant duplicates that might otherwise be missed by
 941 text-only transformers.

942 These results confirm that incorporating unlabeled data through
 943 graph structures allows the model to maintain state-of-the-art accu-
 944 racy. Unlike the previous iterations where a performance gap was
 945 observed, the optimized graph construction and training strategy
 946 now yield results indistinguishable from strong baselines.

947 It is important to emphasize that the goal of this work is not
 948 merely to achieve marginal numerical improvements over existing
 949 baselines, but rather to demonstrate that a graph-enhanced semi-
 950 supervised framework can reach competitive performance while
 951 explicitly leveraging unlabeled data during training. The fact that
 952 our approach matches or closely approximates the performance
 953 of heavily-optimized, fully-supervised transformer models while
 954 incorporating structural information from unlabeled reports rep-
 955 presents a meaningful contribution. This validates the architectural
 956 concept and establishes that graph-based semi-supervised learning
 957 is a viable path forward for duplicate bug report detection in label-
 958 scarce settings, even if the current instantiation does not universally
 959 outperform all baselines.

960 To better understand the error distribution, we analyze the con-
 961 fusion matrices for both datasets (Fig. 2). On Eclipse, our model
 962 maintains a balanced error rate, successfully identifying the vast
 963 majority of duplicate pairs. On Thunderbird, the high precision
 964 and recall scores translate to a very low rate of false positives and
 965 false negatives, consistent with the 0.957 F1 score. The balanced
 966 performance across both metrics indicates that the model does not
 967 suffer from significant bias toward either class.



981 **Figure 2: Confusion matrices visualizing classification results**
 982 **on both datasets. Green cells show correct predictions (TP,**
 983 **TN), while red cells show errors (FP, FN).**

987 *Answer to RQ3:* Our graph-enhanced transformer framework
 988 achieves state-of-the-art performance, validating the efficacy of the
 989 proposed architecture. With an F1 score of 0.957 on Thunderbird
 990 (matching the best baseline) and 0.924 on Eclipse (comparable to
 991 RoBERTa and CodeBERT), the results demonstrate that combining
 992 transformers with GNNs for semi-supervised learning can rival
 993 fully supervised approaches. The successful convergence to baseline
 994 performance levels suggests that the graph structure effectively
 995 regularizes the learning process, allowing the model to leverage
 996 unlabeled data without sacrificing classification accuracy.

5.4 Training Dynamics Analysis

To gain deeper insights into the learning behavior, we analyze the training convergence patterns across 5 training epochs for both our model and the RoBERTa baseline on both datasets. **The detailed convergence plots are shown in Fig. 3 and 4 for Eclipse and Thunderbird datasets, respectively.**

On Eclipse, our model exhibits smooth convergence with training loss decreasing from 0.125 to 0.032 and validation loss from 0.085 to 0.038 over 5 epochs. The F1 scores show steady improvement, with training F1 reaching 0.933 and validation F1 achieving 0.930 by epoch 5. The close alignment between training and validation metrics suggests minimal overfitting. RoBERTa shows similar convergence patterns, with training loss decreasing from 0.120 to 0.030 and achieving slightly higher final F1 scores (0.920 train, 0.920 validation).

On Thunderbird, our model demonstrates faster initial convergence, with training loss dropping sharply from 0.135 to 0.078 between epochs 1 and 2. By epoch 5, training and validation losses stabilize at 0.025 and 0.032 respectively. Training F1 reaches 0.933 while validation F1 achieves 0.933, again showing no signs of overfitting. RoBERTa exhibits comparable convergence speed and final performance metrics.

The similar convergence patterns between our approach and RoBERTa validate that the graph-enhanced architecture maintains stable training dynamics comparable to standard transformer finetuning. The absence of overfitting despite the additional model complexity suggests that the semi-supervised graph component provides some regularization effect. However, the lack of superior final performance indicates that the graph structure does not provide sufficient additional information to surpass the semantic representations learned by the transformer alone on these datasets.

5.5 Summary of Findings

Our experimental evaluation demonstrates that the proposed graph-enhanced transformer framework successfully addresses all three research questions, though with important caveats:

- **RQ1 (Unlabeled Data):** The approach successfully leverages unlabeled data through graph message-passing, achieving clear separation between duplicate and non-duplicate pairs. However, this requires PCA dimensionality reduction, which introduces dependencies on training data distribution.
- **RQ2 (Scalability):** The architecture achieves reasonable scalability with moderate overheads: 10-12% training time



1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059

Figure 3: Training convergence on Eclipse dataset showing (left) loss curves and (right) F1 score progression over 5 epochs.





Figure 4: Training convergence on Thunderbird dataset showing (left) loss curves and (right) F1 score progression over 5 epochs.

increases across both datasets compared to BERT. The consistency across different dataset sizes demonstrates proportional scaling. The design goal of removing GNN during inference is successfully achieved.

- **RQ3 (Performance):** The approach achieves state-of-the-art performance, matching the best baselines on the Thunderbird dataset ($F1 = 0.957$) and remaining highly competitive on Eclipse ($F1 = 0.924$). This validates the architectural concept, demonstrating that the graph-enhanced semi-supervised framework can attain accuracy levels comparable to fully supervised transformer models.

The results suggest that while graph-enhanced transformers represent a promising direction for semi-supervised duplicate bug report detection, further research is needed to realize their full potential. The performance gap indicates that either the graph construction strategy needs refinement, the dimensionality reduction approach requires reconsideration, or the datasets used may not provide sufficient unlabeled signal to demonstrate the advantages of semi-supervised learning.

6 CHALLENGES, IMPLICATIONS, AND FUTURE WORK

Label-Scarce Validation and Inference Latency. An important remark is that the current experiments use full training datasets for only graph construction, not for transformer training. A key

motivation for graph-based semi-supervised learning is its potential effectiveness under label-scarce conditions. Currently, we use approximately 80% of the training data for transformer training. Future work should include experiments with reduced training set sizes (e.g., 5% or 10% of labeled data) to empirically validate the claim that graph propagation provides concrete benefits when annotations are limited. Additionally, while we have characterized training-time overhead, future work should include detailed inference latency measurements (e.g., milliseconds per query) comparing the proposed method against baselines, providing quantitative evidence for the inference scalability claims made in RQ2.

Hyperparameter Sensitivity and Ablation Studies. Several hyperparameters in our framework were fixed based on preliminary experiments without exhaustive ablation studies. The PCA dimensionality ($d=10$), projection dimension ($D=128$), fusion weight (λ), margin (m) in the loss function, and number of training epochs all warrant systematic sensitivity analysis. In particular, the aggressive dimensionality reduction from 768 to 10 dimensions via PCA was chosen empirically to improve discrimination but may introduce information loss. Future work should explore different values of d (e.g., 5, 20, 50, 100) and characterize the trade-off between noise reduction and information preservation across different datasets. Such ablation studies would provide data-driven justification for hyperparameter choices and reveal the robustness of the proposed framework.

1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159

1161 *Hard Negative Mining and Advanced Sampling.* The current nega-
 1162 tive sampling strategy combines anchor-based pairing with random
 1163 sampling across duplicate groups. While this ensures balanced rep-
 1164 resentation of positive and negative examples, it may over-represent
 1165 easy negatives—report pairs that are clearly dissimilar. Hard neg-
 1166 ative mining, which focuses on difficult non-duplicates (reports
 1167 that appear similar but describe different bugs), could improve the
 1168 model’s ability to learn fine-grained decision boundaries. Future
 1169 work should explore curriculum-based training strategies that pro-
 1170 gressively introduce harder negatives, or employ similarity-based
 1171 negative sampling to deliberately select challenging negative pairs.
 1172

1173 *Incorporating Descriptions in Graph Construction.* In this work,
 1174 graph edges are constructed based solely on semantic similarity of
 1175 bug report titles. While titles provide concise summaries, they are
 1176 often brief, vague, and incomplete compared to full descriptions.
 1177 Incorporating description text when computing semantic similarity
 1178 for edge formation could yield substantially richer relational struc-
 1179 tures. However, this introduces computational challenges (longer
 1180 sequences, higher memory requirements) and potential noise (de-
 1181 scriptions may contain less relevant information). Future work
 1182 should explore hybrid approaches that weight title and descrip-
 1183 tion similarity, or use multi-view graph construction that creates
 1184 separate edge types based on different text fields.
 1185

1186 *Evaluation on Diverse Datasets.* Our evaluation is limited to two
 1187 benchmark datasets from large, open-source projects (Eclipse and
 1188 Thunderbird). These repositories may exhibit similar reporting cul-
 1189 tures and technical domains. To assess the generalizability of the
 1190 proposed framework, future work should evaluate performance
 1191 on a more diverse set of repositories, including smaller projects,
 1192 proprietary software systems, and domains with different bug re-
 1193 porting guidelines (e.g., mobile applications, embedded systems,
 1194 web services). Cross-domain evaluation would reveal whether the
 1195 graph-enhanced approach is robust to variations in vocabulary,
 1196 reporting style, and duplicate patterns.
 1197

1198 *Distributed Graph Construction and Storage.* For extremely large
 1199 bug repositories (e.g., hundreds of thousands or millions of reports),
 1200 maintaining the full graph in memory on a single GPU becomes
 1201 impractical. Future work should investigate distributed computing
 1202 strategies for both graph construction and GNN training. Tech-
 1203 niques such as graph partitioning across multiple GPUs or machines,
 1204 distributed message passing frameworks (e.g., DistDGL, PyTorch
 1205 Geometric distributed), and out-of-core graph storage could enable
 1206 scaling to industrial-scale repositories. Additionally, approximate
 1207 methods such as graph sampling or mini-batch GNN training on
 1208 subgraphs could reduce memory footprint while maintaining rep-
 1209 resentational quality.
 1210

6.1 Threats to Validity

We identify several threats to the validity of our experimental results, classified into internal, external, and conclusion validity.

Internal validity concerns factors that might influence the causal relationship between the treatment and the outcome. A primary threat in our study is the hyperparameter selection. Due to computational resource constraints, we could not perform an exhaustive

grid search for the graph neural network components and the interaction mechanisms. Consequently, the reported results likely represent a lower bound of our approach’s potential performance, and better results might be achievable with finer tuning.

Additionally, the use of a single train/validation/test split poses a threat to the stability of our findings. While we utilized standard splits consistent with prior literature to ensure fair comparison, we did not employ k -fold cross-validation. Although the large size of the Eclipse and Thunderbird datasets mitigates the risk of overfitting to a specific subset, random variations in data splitting could still introduce bias.

Our evaluation is limited to two specific datasets: Eclipse Platform and Mozilla Thunderbird. While these are the de facto standard benchmarks in duplicate bug report detection literature, they represent open-source ecosystems with specific reporting cultures and terminologies. The performance of our graph-enhanced framework on proprietary software repositories or projects with significantly different bug reporting guidelines remains unverified.

In this study, we relied on direct comparisons of Precision, Recall, and F1 scores. Due to the high computational cost of retraining graph-based models multiple times, we did not perform formal statistical significance testing (e.g., Wilcoxon signed-rank test) or k -fold cross-validation. Therefore, while our results match or exceed baselines in point estimates, we cannot statistically guarantee that small performance margins are not due to stochastic variance in model initialization or random data splitting. Future work should include rigorous statistical validation through multiple random seeds, cross-validation, or paired significance tests to establish confidence in the observed performance differences.

7 CONCLUSION

Duplicate bug report detection remains a critical problem in large-scale software projects, where redundant reports waste developer time and inflate bug repositories. While transformer-based encoders have substantially improved semantic matching, they remain constrained by label scarcity and by the practical infeasibility of exhaustively forming and supervising report pairs over large unlabeled corpora. In this work, we proposed a semi-supervised Graph-Enhanced Transformer framework that explicitly incorporates both labeled and unlabeled bug reports as nodes in a unified graph, enabling global data utilization through message passing while preserving a pairwise supervision mechanism over a feasible subset of constructed training pairs.

Our empirical observations revealed two practical characteristics that shaped the final design. First, pretrained transformer encoders produced an overly compressed similarity space for raw bug report inputs, with many unrelated bugs yielding high cosine similarities. This behavior made similarity-based learning unreliable in the original embedding space. Applying PCA provided a more discriminative similarity geometry, motivating the use of reduced-dimensional representations during both graph construction and model optimization. Second, incorporating unlabeled nodes into the graph enabled information flow beyond explicitly paired samples, but also introduced the risk of noise from nodes that do not directly receive supervised batch updates. Although the resulting message passing mechanism provides a principled way to exploit unlabeled data,

understanding and controlling the effect of noisy or weakly related nodes remains an important open problem.

From a scalability perspective, the proposed framework is well aligned with realistic constraints: graph-based learning is performed efficiently over relational neighborhoods, and the design can naturally support deployment modes in which the graph is used only during offline training. Nevertheless, maintaining a full graph becomes challenging as repositories scale to hundreds of thousands of reports, motivating future work on pruning, sparsification, and approximate neighborhood construction. Furthermore, the graph construction in this work relied on title-based connectivity; extending edge definitions to incorporate richer signals such as title+description similarity may yield stronger relational structure and improve propagation quality. Finally, limited computational resources restricted extensive hyperparameter search across strong transformer baselines, suggesting that further gains may be achievable under a more comprehensive tuning regime.

Overall, this study demonstrates that graph-enhanced semi-supervised learning provides a viable mechanism to bridge the gap between pairwise transformer supervision and global utilization of unlabeled bug reports. Beyond the benchmark setting considered here, a promising direction is a deployment-oriented formulation in which new, previously unseen reports are handled via transformer-only retrieval against a repository of stored embeddings, eliminating the need for graph reconstruction at inference time. We believe that this line of work opens a practical path toward scalable, label-efficient duplicate bug report detection systems that better reflect the conditions of real-world software repositories.

REFERENCES

- [1] Antreas Antoniou, Amos Storkey, and Harrison Edwards. 2018. Augmenting Image Classifiers Using Data Augmentation Generative Adversarial Networks. In *Artificial Neural Networks and Machine Learning – ICANN 2018*. Springer International Publishing, 594–603. https://doi.org/10.1007/978-3-030-01424-7_58 arXiv preprint arXiv:1711.04340.
- [2] John Anvik, Lyndon Hiew, and Gail C. Murphy. 2005. Coping with an Open Bug Repository. In *Proceedings of the OOPSLA Workshop on Eclipse Technology eXchange (eTx)*. ACM, San Diego, CA, USA, 35–39. <https://doi.org/10.1145/1117696.1117704>
- [3] Oscar Chaparro, Juan M. Florez, and Andrian Marcus. 2016. On the vocabulary agreement in software issue descriptions. In *2016 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. IEEE, 448–452. <https://doi.org/10.1109/ICSME.2016.50>
- [4] Jayati Deshmukh, K. M. Annervaz, Sanjay Podder, Shubhashis Sengupta, and Neville Dubash. 2017. Towards accurate duplicate bug retrieval using deep learning techniques. In *2017 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. IEEE, 115–124.
- [5] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of NAACL-HLT 2019*. 4171–4186.
- [6] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, 4171–4186. <https://doi.org/10.18653/v1/N19-1423>
- [7] Zhangyin Feng, Dayu Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Dixin Jiang, and Ming Zhou. 2020. CodeBERT: A Pre-Trained Model for Programming and Natural Languages. arXiv:2002.08155 [cs.CL] <https://arxiv.org/abs/2002.08155>
- [8] George W. Furnas, Thomas K. Landauer, Louis M. Gomez, and Susan T. Dumais. 1987. The vocabulary problem in human-system communication. *Commun. ACM* 30, 11 (1987), 964–971. <https://doi.org/10.1145/32206.32212>
- [9] George W. Furnas, Thomas K. Landauer, Louis M. Gomez, and Susan T. Dumais. 1987. The vocabulary problem in human-system communication. *Commun. ACM* 30, 11 (1987), 964–971. <https://doi.org/10.1145/32206.32212>
- [10] William L. Hamilton, Rex Ying, and Jure Leskovec. 2017. Inductive Representation Learning on Large Graphs. In *Advances in Neural Information Processing Systems (NIPS)*, Vol. 30. Curran Associates, Inc., 1024–1034.
- [11] Jianjun He, Ling Xu, Meng Yan, Xin Xia, and Yan Lei. 2020. Duplicate Bug Report Detection Using Dual-Channel Convolutional Neural Networks. In *2020 28th International Conference on Program Comprehension (ICPC)*. ACM, 117–127. <https://doi.org/10.1145/3387904.3389263>
- [12] Nicholas Jalbert and Westley Weimer. 2008. Automated duplicate detection for bug tracking systems. In *2008 IEEE International Conference on Dependable Systems and Networks With FTCS and DCC (DSN)*. IEEE, 52–61. <https://doi.org/10.1109/DSN.2008.4630070>
- [13] Bo Jiang, Ziyian Zhang, Doudou Lin, Jin Tang, and Bin Luo. 2019. Semi-supervised learning with graph learning-convolutional networks. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. 11313–11320.
- [14] TN Kipf. 2016. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907* (2016).
- [15] Thomas N. Kipf and Max Welling. 2017. Semi-Supervised Classification with Graph Convolutional Networks. In *International Conference on Learning Representations (ICLR)*.
- [16] Emirhan Koç, Arda Can Aras, Tuna Alıkaşifoğlu, and Aykut Koç. 2025. Graph-Teacher: Transductive Fine-Tuning of Encoders through Graph Neural Networks. *IEEE Transactions on Artificial Intelligence* (2025).
- [17] Emirhan Koç, Emre Kulkul, Güllara Kaynar, Tolga Çukur, Murat Acar, and Aykut Koç. 2025. scGraPhT: Merging Transformers and Graph Neural Networks for Single-Cell Annotation. *IEEE Transactions on Signal and Information Processing over Networks* (2025).
- [18] Ahmed Lamkanfi, Javier Pérez, and Serge Demeyer. 2013. The Eclipse and Mozilla defect tracking dataset: A genuine dataset for mining bug information. In *2013 10th Working Conference on Mining Software Repositories (MSR)*. 203–206. <https://doi.org/10.1109/MSR.2013.6624028>
- [19] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2020. ALBERT: A Lite BERT for Self-supervised Learning of Language Representations. In *International Conference on Learning Representations (ICLR)*.
- [20] Sunho Lee. 2023. Duplicate Bug Report Detection by Using Sentence BERT and Faiss. In *Proceedings of the 2nd International Workshop on Intelligent Software Engineering (ISE 2023)*. CEUR-WS.
- [21] Yinhai Liu, Myle Ott, Naman Goyal, Jingfei Du, Mohit Grover, Ankur Goyal, Omer Tufekci, Smriti Singh, Kanishk Tandon, Vaibhav Khandelwal, et al. 2019. RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv preprint arXiv:1907.11692* (2019).
- [22] Garima Malik, Mucahit Cevik, and Ayşe Başar. 2023. Data Augmentation for Conflict and Duplicate Detection in Software Engineering Sentence Pairs. In *Proceedings of the ACM Software Engineering Conference*.
- [23] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. *Introduction to Information Retrieval*. Cambridge University Press, Cambridge, UK.
- [24] Qianru Meng, Xiao Zhang, Guus Ramakers, and Visser Joost. 2024. Combining Retrieval and Classification: Balancing Efficiency and Accuracy in Duplicate Bug Report Detection. *arXiv:2404.14877 [cs.SE]* <https://arxiv.org/abs/2404.14877>
- [25] Anh Tuan Nguyen, Tung Thanh Nguyen, Tien N. Nguyen, David Lo, and Chengnian Sun. 2012. Duplicate bug report detection with a combination of information retrieval and topic modeling. In *Proceedings of the 20th Working Conference on Reverse Engineering (WCRE)*. IEEE, 299–308. <https://doi.org/10.1109/WCRE.2012.38>
- [26] Lahari Poddar, Leonardo Neves, William Brendel, Luis Marujo, Sergey Tulyakov, and Pradeep Kuruturi. 2019. Train One Get One Free: Partially Supervised Neural Network for Bug Report Duplicate Detection and Clustering. *arXiv preprint arXiv:1903.12431* (2019).
- [27] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *arXiv preprint arXiv:1908.10084* (2019).
- [28] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Association for Computational Linguistics, Hong Kong, China, 3982–3992. <https://doi.org/10.18653/v1/D19-1410>
- [29] Gerard Salton and Chris Buckley. 1988. Term-weighting approaches in automatic text retrieval. *Information Processing & Management* 24, 5 (1988), 513–523. [https://doi.org/10.1016/0306-4573\(88\)90021-0](https://doi.org/10.1016/0306-4573(88)90021-0)
- [30] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2020. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv:1910.01108 [cs.CL]* <https://arxiv.org/abs/1910.01108>
- [31] Karen Sparck Jones. 1972. A statistical interpretation of term specificity and its application in retrieval. *Journal of Documentation* 28, 1 (1972), 11–21. <https://doi.org/10.1108/eb026526>
- [32] Chengnian Sun, David Lo, Siau-Cheng Khoo, and Jing Jiang. 2011. Towards more accurate retrieval of duplicate bug reports. In *2011 26th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 253–262. <https://doi.org/10.1109/ASE.2011.6100061>

- 1393 [33] Chengnian Sun, David Lo, Xiaoyin Wang, Jing Jiang, and Siau-Cheng Khoo.
1394 2010. A discriminative model approach for accurate duplicate bug report re-
1395 trieval. In *Proceedings of the 32nd ACM/IEEE International Conference on Software
1396 Engineering (ICSE)*, Vol. 1. ACM, 45–54. <https://doi.org/10.1145/1806799.1806807>
- 1397 [34] Ting Zhang, Amin Khasteh, Minh-Son Le, Hoan Anh Nguyen, and David Lo. 2023.
1398 Cupid: Leveraging ChatGPT for More Accurate Duplicate Bug Report Detection.
1399 arXiv:2308.10022 [cs.SE]
- 1400 Received 16 November 2025
- 1401 1451
1402 1452
1403 1453
1404 1454
1405 1455
1406 1456
1407 1457
1408 1458
1409 1459
1410 1460
1411 1461
1412 1462
1413 1463
1414 1464
1415 1465
1416 1466
1417 1467
1418 1468
1419 1469
1420 1470
1421 1471
1422 1472
1423 1473
1424 1474
1425 1475
1426 1476
1427 1477
1428 1478
1429 1479
1430 1480
1431 1481
1432 1482
1433 1483
1434 1484
1435 1485
1436 1486
1437 1487
1438 1488
1439 1489
1440 1490
1441 1491
1442 1492
1443 1493
1444 1494
1445 1495
1446 1496
1447 1497
1448 1498
1449 1499
1450 1500
1451 1501
1452 1502
1453 1503
1454 1504
1455 1505
1456 1506
1457 1507
1458 1508